

Likelihood-Ranked Synthetic Selection in Self-Consuming Diffusion: Minority Extinction Ordering on Imbalanced Mixtures

Muhammad Alamzeb
Independent researcher
shayankhanmahar@gmail.com

Abstract

Self-consuming training loops—retraining generative models on their own outputs—can induce *model collapse*, including loss of rare modes. Prior work notes that bias toward high-quality synthetic samples worsens the quality–diversity tradeoff, and that task-level verification can stabilize labeled synthetic data. We study the most common *unlabeled* selection score for diffusion models: a denoising ELBO / likelihood proxy. On imbalanced Gaussian mixtures with a tiny DDPM, under matched real-data mixing and matched synthetic budget, first-generation minority-mass retention obeys

$$\text{bottom-}k > \text{random-}k > \text{top-}k$$

(means 0.57/0.36/0.17 over eight seed–imbalance pairs; order holds on 5/5 seeds at imbalance $\rho=5$). Paired Wilcoxon tests vs. random- k give $p \approx 0.042$ for both top and bottom contrasts. Likelihood top- k also concentrates majority mass relative to random- k . A natural “false stability” hypothesis—that top- k improves sliced W_2 while minorities die—is falsified. The ordering transfers to 8D GMM and to Digits in PCA space with nearest class-mean occupancy (3/3 seeds each). We release configs, seeds, logs, and analysis code.

1 Introduction

As synthetic data proliferates, generative models are increasingly trained on mixtures of real and model-generated samples. Iterated self-consumption can cause *model collapse*: progressive loss of diversity and forgetting of rare modes [Shumailov et al., 2023, Alemohammad et al., 2023, Gerstgrasser et al., 2024]. Mitigations include accumulating real data [Gerstgrasser et al., 2024], injecting fresh real samples, and *filtering* synthetic data [Feng et al., 2024, Cai et al., 2025].

A practical question remains underspecified: when the only available score is the generator’s own likelihood (or denoising loss), how should one select synthetic samples? Feng et al. [2024] show that *task verifiers* for labeled synthetic answers can prevent collapse. Cai et al. [2025] filter “unrealistic” diffusion samples via latent probes and report aggregate FID/precision/recall. Alemohammad et al. [2023] observe that sampling bias toward high-quality generations drives a quality–diversity tradeoff. What is missing is a **matched-budget, mode-exact** causal contrast among likelihood ranks.

Contributions.

1. A controlled self-consuming DDPM protocol on imbalanced GMMs with exact mode masses.
2. Evidence for a **likelihood-rank ordering** of minority retention under matched k and matched real mix ratio.

3. An experimental **falsification** of W_2 -based false stability for likelihood top- k .
4. Transfer checks: 8D GMM and Digits-PCA with nearest class-mean occupancy.

2 Related work

Model collapse. Shumailov et al. [2023] document recursive training collapse and tail forgetting. Gerstgrasser et al. [2024] show accumulating real+synthetic data can avoid collapse. Alemohammad et al. [2023] (*MAD*) highlight quality-biased sampling.

Verification and filtering. Feng et al. [2024] argue verification of synthesized *labels* prevents collapse. Cai et al. [2025] (*LSF*) filter synthetic images using latent-space realism scores. Neither isolates unlabeled ELBO top- k vs. random- k with mode-grounded metrics on continuous mixtures.

Temperature / diversity knobs. Inference-time temperature methods [Xu et al., 2025, Li et al., 2026] reshape sampling diversity; we instead intervene on *training-set selection* inside self-consuming loops.

3 Method

3.1 Self-consuming loop

Let $f^{(g)}$ be a DDPM trained on $\mathcal{D}^{(g)}$. Sample synthetic set $\hat{\mathcal{D}}^{(g+1)}$. Build the next train set by mixing a fixed real buffer $\mathcal{D}_{\text{real}}$ (fraction α) with selected synthetics.

3.2 Likelihood proxy and policies

For each synthetic x , define score $s(x)$ as mean denoising MSE over a grid of times t (ELBO proxy; **lower** = higher likelihood). Policies at matched synthetic count: **top- k** (lowest s), **random- k** , **bottom- k** (highest s), plus **mix** / **replace** references.

3.3 Metrics

Nearest-mean mode assignment on known centers (GMM means or Digits empirical class means in PCA space). Report minority-mean mass, majority mass, and sliced W_2 to held-out real samples.

4 Experiments

Setup. $K=4$ modes, $\rho \in \{5, 10\}$, tiny MLP denoiser, $G=5$, seeds $\{0 \dots 4\}$ (seeds 3–4 replicate $\rho=5$), $\alpha=0.5$. Logs: `experiments/logs/w2_fs.jsonl` ($n=228$).

At $\rho=5$, generation 1, all five seeds satisfy bottom > rand > top. Wilcoxon signed-rank vs. random- k ($n=8$): top–rand $\Delta=-0.188$ ($p \approx 0.042$); bottom–rand $\Delta=+0.212$ ($p \approx 0.042$). Top- k does **not** improve sliced W_2 (mean $\Delta W_2 > 0$).

Real-mix sensitivity. An α -sweep (`w2_fs_alpha.jsonl`) shows the top- k penalty is clearest at $\alpha \leq 0.5$; at $\alpha=0.75$ it weakens vs. random while bottom- k remains highest on average.

Table 1: Minority-mass retention $m^{(1)}/m^{(0)}$ on the principal GMM matrix.

Policy	Mean	Std	n
bottom_k	0.572	0.289	8
rand_k	0.360	0.173	8
mix	0.328	0.100	8
top_k	0.172	0.078	8




Figure 1: Left: retention by policy. Right: minority mass vs. generation at $\rho=5$.

Transfers. An 8D GMM transfer (`w2_fs_hd.jsonl`) reproduces the ordering on 3/3 seeds (retention means 0.55/0.27/0.08). Digits in PCA-8 with nearest class-mean occupancy (`w2_fs_digits_pca.jsonl`) holds on 3/3 seeds (retention means 1.49/1.11/0.68). An earlier Digits CVAE+classifier probe is inconclusive and unused for claims.

Selection-pool size. Under $k = \max(n_{\text{syn}}, \lfloor k_{\text{frac}} N \rfloor)$, $k_{\text{frac}} \in \{0.25, 0.5\}$ are operationally equivalent when n_{syn} binds; at $k_{\text{frac}}=0.75$ top- k remains worst.

5 Analysis

Likelihood top- k preferentially retains synthetic points the current model already explains well—typically majority-mode samples—reducing minority inflow into $\mathcal{D}^{(g+1)}$. Bottom- k does the opposite relative to random. This is consistent with MAD’s quality-bias narrative but makes the score-rank interventional and mode-exact. Aggregate W_2 is a poor monitor for this failure: it moves with top- k in the *wrong* direction for a “quality win.”

6 Limitations

Toy 2D/8D GMMs and Digits-PCA with tiny networks; not ImageNet-scale. Digits CVAE transfer was inconclusive; Digits-PCA uses embeddings, not raw pixels. Effect strongest at moderate imbalance and early generations; high real mix ($\alpha=0.75$) weakens the top- k vs. random gap. Novelty is **L2 empirical** relative to MAD/LSF/Feng; we do not claim a new SOTA training method on FID.

7 Conclusion

Unlabeled likelihood ranking of synthetic data in self-consuming diffusion induces a stable minority-retention order (bottom- $k > \text{random-}k > \text{top-}k$) under matched budgets on imbalanced mixtures, while failing to deliver W_2 improvements. Preferring “most likely” synthetics is not a free lunch for minority survival.

References

- Sina Alemohammad, Josue Casco-Rodriguez, Lorenzo Luzi, Ahmed Imtiaz Humayun, Hossein Babaei, Daniel LeJeune, Ali Siahkoohi, and Richard G. Baraniuk. Self-consuming generative models go MAD, 2023.
- Zhongteng Cai, Yaxuan Wang, Yang Liu, and Xueru Zhang. Stabilizing self-consuming diffusion models with latent space filtering, 2025.
- Yunzhen Feng, Elvis Dohmatob, Pu Yang, Francois Charton, and Julia Kempe. Beyond model collapse: Scaling up with synthesized data requires verification, 2024.
- Matthias Gerstgrasser, Rylan Schaeffer, Apratim Dey, Rafael Rafailov, Dhruv Pai, Henry Sleight, John Hughes, Tomasz Korbak, Rajashree Agrawal, Andrey Gromov, Daniel A. Roberts, Diyi Yang, David Donoho, and Sanmi Koyejo. Is model collapse inevitable? breaking the curse of recursion by accumulating real and synthetic data, 2024.
- Peizhuo Li, Emre Aksan, Alexandru-Eugen Ichim, Thabo Beeler, and Olga Sorkine-Hornung. Diversify diffusion with temperature sampling and variance-corrective time shifting, 2026.
- Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Yarin Gal, Nicolas Papernot, and Ross Anderson. The curse of recursion: Training on generated data makes models forget, 2023.
- Yanbo Xu, Yu Wu, Sungjae Park, Zhizhou Zhou, and Shubham Tulsiani. Temporal score rescaling for temperature sampling in diffusion and flow models, 2025.

A Reproducibility

Principal command (Windows PowerShell): `scripts/reproduce_w2.ps1`. All numerical tables are produced from logged JSONL via analysis scripts (no hand-edited metrics). Experiment IDs: `w2_fs`, `w2_fs_hd`, `w2_fs_digits_pca` (supported); `w2_fs_digits` (CVAE; inconclusive). Code and logs accompany this manuscript.